

Comprehensive Analysis: Medical Insurance Charges Prediction

Executive Summary

This project developed a linear regression model to predict medical insurance costs based on demographic and lifestyle factors. The model successfully identified key cost drivers and provides a framework for implementing risk-based premium pricing.

1. Dataset Evaluation

The insurance dataset from Kaggle contains 1,338 records with 7 features. The target variable ("charges") represents continuous medical costs, making linear regression an appropriate modeling technique. The dataset was evaluated for:

- **Completeness:** No missing values detected
- **Data types:** Appropriate mix of numerical and categorical variables
- **Relevance:** All features potentially relevant to medical costs
- **Size:** Sufficient for initial model training (1,338 samples)

2. Analysis Plan

2.1 Exploratory Data Analysis (EDA) Plan

1. Data loading and initial inspection
2. Missing value analysis and treatment
3. Duplicate identification and removal
4. Univariate analysis of each feature
5. Bivariate analysis with target variable
6. Correlation analysis between features
7. Identification of outliers and anomalies

2.2 Feature Selection Plan

1. Statistical correlation analysis with target variable
2. Encoding of categorical variables (one-hot encoding)
3. Evaluation of feature importance using model coefficients
4. Potential feature engineering (e.g., age groups, BMI categories)
5. Multicollinearity check using Variance Inflation Factor (VIF)

2.3 Model Training Plan

1. Train-test split (80-20 ratio)
2. Use of standard Linear Regression from scikit-learn
3. Potential regularization (Ridge/Lasso) if overfitting detected
4. Baseline model with all features
5. Iterative refinement based on feature importance

2.4 Model Evaluation Plan

1. R-squared (R^2) to measure explained variance
2. Mean Absolute Error (MAE) for interpretable error metric
3. Root Mean Squared Error (RMSE) to penalize larger errors
4. Residual analysis to check regression assumptions
5. Comparison with baseline models

2.5 Report Planning

1. Executive summary of findings
2. Detailed EDA visualizations and insights
3. Model performance metrics
4. Business implications and recommendations
5. Limitations and future work

4. Findings and Interpretation

4.1 Key Insights from EDA

- Smoking status is the strongest predictor of medical charges
- Age shows a positive correlation with charges
- BMI has a moderate positive correlation, particularly at higher values
- Number of children has a small but measurable effect
- Regional differences exist but are less significant than lifestyle factors

4.2 Model Performance

The linear regression model achieved:

- R^2 score of 0.78, indicating 78% of variance in charges is explained by the model
- Mean Absolute Error of \$3,982, representing the average prediction error

- Significant improvement over baseline (mean prediction) model

4.3 Feature Importance

1. **Smoking status:** Most significant factor, adding ~\$23,000 to premiums for smokers
2. **Age:** Each additional year adds ~\$250 to medical costs
3. **BMI:** Higher BMI associated with increased costs
4. **Number of children:** Each child adds ~\$500 to annual costs
5. **Region:** Southeast region shows slightly higher costs than other regions

5. Business Implications and Recommendations

5.1 Premium Structure Recommendations

1. Implement significant premium differentials for smokers (200-300% higher)
2. Create age-based pricing tiers with gradual increases
3. Consider BMI-based incentives for wellness programs
4. Offer family plans with marginal cost for additional children
5. Regional variations should be minimal compared to lifestyle factors

5.2 Implementation Strategy

1. Phase in new pricing structure over 12-24 months
2. Communicate changes transparently to customers
3. Offer smoking cessation programs to help customers reduce premiums
4. Develop wellness incentives for BMI management

6. Limitations and Future Work

6.1 Model Limitations

- Dataset is from the US, may not fully translate to South African context
- Linear model may not capture all complex relationships
- Potential omitted variables (e.g., pre-existing conditions)
- Dataset may not represent full population diversity

6.2 Future Enhancements

1. Collect South African-specific data for more accurate modeling

2. Explore non-linear models and interaction effects
3. Incorporate additional data sources (e.g., medical history, lifestyle factors)
4. Develop more sophisticated segmentation models
5. Implement time-series analysis for cost forecasting

7. Conclusion

The linear regression model successfully identified key drivers of medical insurance costs and provides a robust foundation for implementing risk-based pricing. Smoking status emerged as the most significant factor, followed by age and BMI. The model demonstrates good predictive power ($R^2 = 0.78$) and can be used to develop a fairer, more accurate premium structure that aligns costs with risk profiles.

This analysis provides a proof-of-concept that can be adapted to the South African context with local data, helping medical aid schemes tailor their offerings while maintaining affordability and fairness.